

k-means Clustering

TOP: Data Clustering 076/091

Instructor: Sayan Bandyapadhyay
Portland State University

Outline

1 The k -means Algorithm

2 Quality Analysis of k -means

Real Points

Suppose the set of points X are from \mathbb{R}^d

- A natural center of points is the average point or mean

$$\mu = \frac{1}{|S|} \cdot \sum_{x \in S} x$$

- Here the sum is coordinate-wise total:

$$(1, 3) + (2, 5) = (3, 8)$$

Real Points

Suppose the set of points X are from \mathbb{R}^d

- A natural center of points is the average point or mean

$$\mu = \frac{1}{|S|} \cdot \sum_{x \in S} x$$

- Here the sum is coordinate-wise total:
 $(1, 3) + (2, 5) = (3, 8)$
- This is the basis of the k -means algorithm
- Proposed by Lloyd in 1957, published in 1982
- Also by Max in 1960

Real Points

Suppose the set of points X are from \mathbb{R}^d

- A natural center of points is the average point or mean

$$\mu = \frac{1}{|S|} \cdot \sum_{x \in S} x$$

- Here the sum is coordinate-wise total:

$$(1, 3) + (2, 5) = (3, 8)$$

- This is the basis of the k -means algorithm
- Proposed by Lloyd in 1957, published in 1982
- Also by Max in 1960

Euclidean distance of x and y , $\|x - y\| = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$

The k -means Algorithm (Lloyd-Max)

Algorithm k -means

Require: Set of points X

- 1: Start with centers c_1, \dots, c_k chosen arbitrarily from X
 - 2: **repeat**
 - 3: **for** each point $x_i \in X$ **do**
 - 4: Assign x_i to cluster C_j that minimizes $\|x_i - c_j\|$
 - 5: **end for**
 - 6: **for** each cluster C_j **do**
 - 7: $c_j \leftarrow \frac{1}{|C_j|} \cdot \sum_{x_i \in C_j} x_i$
 - 8: **end for**
 - 9: **until** cluster centers do not change
-

Time Complexity of k -means

Algorithm k -means

Require: Set of points X

- 1: Start with centers c_1, \dots, c_k chosen arbitrarily from X
 - 2: **repeat**
 - 3: **for** each point $x_i \in X$ **do**
 - 4: Assign x_i to cluster C_j that minimizes $\|x_i - c_j\|$
 - 5: **end for**
 - 6: **for** each cluster C_j **do**
 - 7: $c_j \leftarrow \frac{1}{|C_j|} \cdot \sum_{x_i \in C_j} x_i$
 - 8: **end for**
 - 9: **until** cluster centers do not change
-

Time Complexity of k -means

Algorithm k -means

Require: Set of points X

- 1: Start with centers c_1, \dots, c_k chosen arbitrarily from X
 - 2: **repeat**
 - 3: **for** each point $x_i \in X$ **do**
 - 4: Assign x_i to cluster C_j that minimizes $\|x_i - c_j\|$
 - 5: **end for**
 - 6: **for** each cluster C_j **do**
 - 7: $c_j \leftarrow \frac{1}{|C_j|} \cdot \sum_{x_i \in C_j} x_i$
 - 8: **end for**
 - 9: **until** cluster centers do not change
-

- Again we need a “rate of cost decrease” type argument as for k -median
- What is a suitable cost function that the mean minimizes?

A Suitable Cost Function

- For what function $g(., .)$, $\text{mean}(\mathcal{S})$ minimizes $\sum_{x \in \mathcal{S}} g(x, c)$ over all c ?

A Suitable Cost Function

- For what function $g(., .)$, $\text{mean}(\mathcal{S})$ minimizes $\sum_{x \in \mathcal{S}} g(x, c)$ over all c ?
- Such g is called Bregman divergence that encompasses many functions

A Suitable Cost Function

- For what function $g(., .)$, $\text{mean}(\mathcal{S})$ minimizes $\sum_{x \in \mathcal{S}} g(x, c)$ over all c ?
- Such g is called Bregman divergence that encompasses many functions
- One such g is squared Euclidean distance

$$g(x, c) = ||x - c||^2$$

A Suitable Cost Function

- For what function $g(\cdot, \cdot)$, $\text{mean}(\mathcal{S})$ minimizes $\sum_{x \in \mathcal{S}} g(x, c)$ over all c ?
- Such g is called Bregman divergence that encompasses many functions
- One such g is squared Euclidean distance

$$g(x, c) = ||x - c||^2$$

- This leads to our k -means clustering problem for real points with Euclidean distance

k-means clustering

Given a set X of n points in the metric space (\mathcal{U}, d)

- Find a set C of k points (cluster centers) in \mathcal{U} that minimizes,

$$\text{cost}(C) = \sum_{p \in X} d(p, \text{NearestCenter}(p))^2$$

Euclidean k -means clustering

Given a set X of n points in \mathbb{R}^d

- Find a set C of k points (cluster centers) in \mathbb{R}^d that minimizes,

$$\text{cost}(C) = \sum_{p \in X} \|p - \text{NearestCenter}(p)\|^2$$

Time Complexity of Lloyd's Algorithm

- M_1, M_2, \dots, M_ℓ are the sets of means computed over ℓ iterations

To show: $\text{cost}(M_\ell) < \text{cost}(M_{\ell-1}) < \text{cost}(M_{\ell-2}) < \dots < \text{cost}(M_1)$

Time Complexity of Lloyd's Algorithm

- M_1, M_2, \dots, M_ℓ are the sets of means computed over ℓ iterations

To show: $\text{cost}(M_\ell) < \text{cost}(M_{\ell-1}) < \text{cost}(M_{\ell-2}) < \dots < \text{cost}(M_1)$

- In every iteration, means are picked as centers of the clusters
- A mean minimizes the sum-of-squares cost function
- So, k -means cost also decreases for the new set of centers

Time Complexity

- In every iteration, cost decreases
- The algorithm never cycles – The same set of centers never comes back
- Number of iterations is bounded by the number of distinct sets of means

Time Complexity

- In every iteration, cost decreases
- The algorithm never cycles – The same set of centers never comes back
- Number of iterations is bounded by the number of distinct sets of means
- 2^n subsets: 2^n distinct means; $(2^n)^k$ distinct sets of means

Time Complexity


- In every iteration, cost decreases
- The algorithm never cycles – The same set of centers never comes back
- Number of iterations is bounded by the number of distinct sets of means
- 2^n subsets: 2^n distinct means; $(2^n)^k$ distinct sets of means
- Lloyd's algorithm always terminates
- In practice, it is very fast
- One can also terminate the algorithm after a few iterations

Outline

1 The k -means Algorithm

2 Quality Analysis of k -means

Analysis of Quality



Initialization/Seeding is the key

Initialization

- Does random initialization help?

Initialization

- Does random initialization help? No! We still can get nearby centers

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?
- Sensitive to outliers

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?
- Sensitive to outliers
- We need something in between

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?
- Sensitive to outliers
- We need something in between
- we should pick far away points only if there are many points in the vicinity

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?
- Sensitive to outliers
- We need something in between
- we should pick far away points only if there are many points in the vicinity
- We should not pick an outlier as a center

Initialization

- Does random initialization help? No! We still can get nearby centers
- We need well-separated centers – can we use Greedy 2 - Furthest point algorithm for k -center?
- Sensitive to outliers
- We need something in between
- we should pick far away points only if there are many points in the vicinity
- We should not pick an outlier as a center

This leads to a new seeding algorithm!

Non-uniformly Random Seeding

- Start with a uniformly random center

Non-uniformly Random Seeding

- Start with a uniformly random center
- Next center is chosen from a distribution biased towards far away points

Non-uniformly Random Seeding

- Start with a uniformly random center
- Next center is chosen from a distribution biased towards far away points
- Cost of a point x_i , $\text{cost}(x_i, C) = \min_{c \in C} \|x_i - c\|^2$

Non-uniformly Random Seeding

- Start with a uniformly random center
- Next center is chosen from a distribution biased towards far away points
- Cost of a point x_i , $\text{cost}(x_i, C) = \min_{c \in C} \|x_i - c\|^2$
- Define the probability $p_i = \text{cost}(x_i, C) / \text{cost}(C)$

The k -means++ Algorithm

Algorithm k -means++

Require: Set of points X , parameter k

- 1: Select c_1 randomly from X
 - 2: $C \leftarrow \{c_1\}$
 - 3: **while** ($|C| \neq k$) **do**
 - 4: **for** each $i = 1$ to n **do**
 - 5: $\text{cost}(x_i, C) \leftarrow \min_{c \in C} \|x_i - c\|^2$
 - 6: **end for**
 - 7: $\text{cost}(C) \leftarrow \sum_{i=1}^n \text{cost}(x_i, C)$
 - 8: Sample a random point $y \in X$, selecting each x_i w.p.
 $p_i = \text{cost}(x_i, C) / \text{cost}(C)$
 - 9: $C \leftarrow C \cup \{y\}$
 - 10: **end while**
 - 11: Invoke Lloyd's algorithm with C as the seed
-

Analysis of k -means++

- Time complexity: $O(nk)$ + Lloyd's

Analysis of k -means++

- Time complexity: $O(nk)$ + Lloyd's
- Approximation factor: $O(\log k)$

Analysis of k -means++

- Time complexity: $O(nk)$ + Lloyd's
- Approximation factor: $O(\log k)$

Compare this with p -swap Local search

- $O(n^{p+1} k^{p+1} \log n)$ time, but $9 + (1/p)$ -approximation
- Works in general metric space (even for non-numerical data)